


Physics-Informed Neural Networks (PINNs) Basics

By,
Faiq Shahbaz



Problem Setting: What Are We Solving?

- We want to solve a PDE of the general form:

$$u_t + \mathcal{N}[u] = 0$$

- unknown function (e.g. velocity, temperature).
- t: time, x: space.
- $\mathcal{N}[u]$: some (possibly nonlinear) expression involving u and its derivatives (for example, $uu_x - vu_{xx}$).
- We also have:
 - Initial condition (IC): $u(0, x) = u_0(x)$
 - Boundary conditions (BC): $u(t, x) = g(t, x)$ on the boundary of the spatial domain.

Core Idea of a PINN

- Use a neural network to approximate the solution:

$$u_{\theta}(t, x) \approx u(t, x)$$

- Train this neural network to satisfy:

- Data (initial and boundary conditions).
- The PDE (physics) inside the domain.

- Define a total loss function:

$$\text{Loss}(\theta) = \text{Loss}_{\text{data}}(\theta) + \text{Loss}_{\text{physics}}(\theta)$$

- $\text{Loss}_{\text{data}}$: mismatch at IC and BC points.
- $\text{Loss}_{\text{physics}}$: mismatch of the PDE residual (how much we violate the equation).

What Is Automatic Differentiation (AD)

- Automatic differentiation (AD) is how frameworks (PyTorch, TensorFlow, etc.) compute derivatives automatically.
- When you compute a function, the library builds a **computation graph**:
 - Nodes = intermediate values (results of operations).
 - Edges = operations (add, multiply, tanh, etc.).
- Each operation has a known derivative.
- AD runs a **backward pass** over this graph using the **chain rule** to compute derivatives:
 - With respect to parameters (weights) → standard backprop.
 - With respect to inputs (like t and x) → needed for u_t, u_x, u_{xx} , etc.

Computation Graph Example

- Consider:

$$f(x) = x^2 + 3x$$

- Break into intermediate nodes:

- Node 0 (input): x
- Node 1: $a = x^2$
- Node 2: $b = 3x$
- Node 3 (output): $c = a + b$ (so $f(x) = c$)

- Local derivatives:

$$\frac{\partial a}{\partial x} = 2x, \frac{\partial b}{\partial x} = 3, \frac{\partial c}{\partial a} = 1, \frac{\partial c}{\partial b} = 1$$

- Total derivative via chain rule:

$$\frac{dc}{dx} = \frac{\partial c}{\partial a} \frac{\partial a}{\partial x} + \frac{\partial c}{\partial b} \frac{\partial b}{\partial x}$$

$$\frac{dc}{dx} = 1 \cdot 2x + 1 \cdot 3 = 2x + 3$$

Tiny Neural Network Example

- Define a very small neural network:

$$z = w_1 t + w_2 x + b$$

$$u = \tanh(z)$$

- Here:

- Inputs: t and x
- Parameters: w_1, w_2, b
- Output: $u(t, x)$
- Local derivatives of the li

$$\frac{\partial z}{\partial t} = w_1, \frac{\partial z}{\partial x} = w_2, \frac{\partial z}{\partial w_1} = t, \frac{\partial z}{\partial w_2} = x, \frac{\partial z}{\partial b} = 1$$

Activation Derivative

- We have:
- Known calculus result:

$$u = \tanh(z)$$

$$\frac{d}{dz} \tanh(z) = 1 - \tanh^2(z)$$

- Therefore:

$$\frac{\partial u}{\partial z} = 1 - \tanh^2(z)$$

- Since $u = \tanh(z)$, we can write:

$$\tanh^2(z) = u^2$$

- So:

$$\frac{\partial u}{\partial z} = 1 - u^2$$

Derivatives from the Tiny Network

- Using chain rule for input derivatives:

$$\frac{\partial u}{\partial t} = \frac{\partial u}{\partial z} \frac{\partial z}{\partial t} = (1 - u^2) w_1$$

$$\frac{\partial u}{\partial x} = \frac{\partial u}{\partial z} \frac{\partial z}{\partial x} = (1 - u^2) w_2$$

- Derivatives w.r.t. parameters (like in backprop):

$$\frac{\partial u}{\partial w_1} = \frac{\partial u}{\partial z} \frac{\partial z}{\partial w_1} = (1 - u^2) t$$

$$\frac{\partial u}{\partial w_2} = (1 - u^2) x$$

$$\frac{\partial u}{\partial b} = (1 - u^2) 1 = 1 - u^2$$

- AD in frameworks computes these automatically from the computation graph.

Designing the Neural Network in a PINN

- We design a neural network NN_θ to approximate the solution:

$$u_\theta(t, x) = \text{NN}_\theta(t, x)$$

- Inputs:
 - All independent variables, e.g. (t, x) for 1D time-dependent problems.
- Outputs:
 - The unknown field(s), e.g. scalar u or vector (u, v, p) for systems.
- Typical architecture (for PINNs):
 - 4–10 hidden layers.
 - 20–100 neurons per layer.
 - Smooth activation functions: commonly \tanh or \sin .

Types of Training Points

- **Initial condition (IC) points:**
 - Points on $t = 0$, with known values:
 - $u(0, x_i) = u_0(x_i)$
- **Boundary condition (BC) points:**
 - Points on spatial boundaries, e.g. $x = -1$ and $x = 1$.
 - Values given by the boundary condition function $g(t, x)$.
- **Collocation (interior) points:**
 - Randomly sampled points inside the space–time domain.
 - No target u value is given.
 - Used to enforce that the PDE residual is close to zero.

Example PDE: Burgers' Equation

- Burgers' equation:

$$u_t + uu_x - \nu u_{xx} = 0$$

- Domain:

$$x \in [-1,1], t \in [0,1]$$

- Viscosity:

$$\nu = \frac{0.01}{\pi}$$

- Initial condition:

$$u(0, x) = -\sin(\pi x)$$

- Boundary conditions:

$$u(t, 1) = 0, \quad u(t, -1) = 0$$

PINN Approximation for Burgers

- Define neural network solution:

$$u_\theta(t, x) = \text{NN}_\theta(t, x)$$

- At any point (t, x) :

- Network outputs an approximation of $u(t, x)$.

- To enforce the PDE, we need:

$$u_t(t, x) = \frac{\partial u_\theta}{\partial t}(t, x)$$

$$u_x(t, x) = \frac{\partial u_\theta}{\partial x}(t, x)$$

$$u_{xx}(t, x) = \frac{\partial^2 u_\theta}{\partial x^2}(t, x)$$

- These derivatives are computed using automatic differentiation.

PDE Residual Using AD

- At collocation point (t, x):

- Compute $u = u_\theta(t, x)$.
- Use AD to compute:

$$u_t = \frac{\partial u_\theta}{\partial t}(t, x)$$

$$u_x = \frac{\partial u_\theta}{\partial x}(t, x)$$

$$u_{xx} = \frac{\partial^2 u_\theta}{\partial x^2}(t, x)$$

- Define the PDE residual:

$$f_\theta(t, x) = u_t + u_x - vu_{xx}$$

- If the PDE is satisfied exactly, then:

$$f_\theta(t, x) = 0$$

Loss Terms for Burgers' PINN

- **Initial condition loss** (Loss_{IC}):

- For IC points

$$\text{Loss}_{\text{IC}} = \frac{1}{N_0} \sum_{i=1}^{N_0} (u_\theta(0, x_i) - (-\sin(\pi x_i)))^2$$

- **Boundary condition loss** (Loss_{BC}):

- For boundary points

$$\text{Loss}_{\text{BC}} = \frac{1}{N_b} \sum_{j=1}^{N_b} (u_\theta(t_j, -1))^2 + \frac{1}{N_b} \sum_{j=1}^{N_b} (u_\theta(t_j, 1))^2$$

- **PDE residual loss** (Loss_{PDE}):

- For collocation points (t_k, x_k) :

$$\text{Loss}_{\text{PDE}} = \frac{1}{N_f} \sum_{k=1}^{N_f} f_\theta(t_k, x_k)^2$$

- **Total loss**:

- $\text{Loss}(\theta) = \text{Loss}_{\text{IC}} + \text{Loss}_{\text{BC}} + \text{Loss}_{\text{PDE}}$

Training Pipeline (Algorithm View)

- **Step 1:** Choose neural network architecture.
 - Input size = 2 (t, x), output size = 1 (u).
- **Step 2:** Initialize network parameters θ randomly.
- **Step 3:** Sample three sets of points:
 - IC points: $(0, x_i)$ with $u_0(x_i) = -\sin(\pi x_i)$.
 - BC points: $(t_j, -1)$ and $(t_j, 1)$ with value 0.
 - Collocation points: (t_k, x_k) randomly in $[0, 1] \times [-1, 1]$.
- **Step 4:** For each training iteration:
 - Compute u_θ at IC and BC points and evaluate Loss_{IC} and Loss_{BC} .
 - At collocation points, use AD to compute u_t, u_x, u_{xx} and then $f_\theta(t_k, x_k)$.
 - Compute Loss_{PDE} from residuals.
 - Sum to get total $\text{Loss}(\theta)$.
 - Use backpropagation (AD) to compute $\nabla_\theta \text{Loss}$.
 - Update θ with an optimizer (e.g. gradient descent, Adam).