


# **Solving Burger's Equation with Physics-Informed Neural Networks (PINNs)**

By,  
**Faiq Shahbaz**



# Motivation, intuition, and where PINNs fit

---

- **Problem class:** Solve/identify PDE solutions when labeled data are scarce but governing laws are known.
- **Key idea:** Replace large labeled datasets with the **PDE residual** and known **IC/BC** as supervision.
- **Conceptual loop:** Sample points → predict  $u_\theta$  → compute derivatives via autodiff → form residual → minimize residual + constraint penalties.
- **Why this matters:**
  - Works with **little or no data**; integrates physics priors; differentiable solution useful for gradients/sensitivities.
  - Natural path to **inverse problems** (identify parameters/fields from sparse measurements).
- **This slide:** from PDE and function spaces → continuous PINN loss → discrete Monte-Carlo estimator → autodiff & normalization → optimization → diagnostics → benchmark & error analysis → extensions.

PINN: minimize physics residual and constraints

$$\min_{\theta} \left( \|\mathcal{N}[u_\theta]\|_{L^2(\Omega)}^2 + \lambda_{ic} \|u_\theta(\cdot, 0) - u_0\|_{L^2(-1,1)}^2 + \lambda_{bc} \|u_\theta(\pm 1, \cdot)\|_{L^2(0,1)}^2 \right).$$

# PDE, domain, physical meaning, non-dimensionalization

---

- **PDE:** Viscous Burgers models 1D convection–diffusion with nonlinear advection and linear diffusion.
- **Domain:**  $\Omega = [-1,1] \times [0,1]$  (*space*  $\times$  *time*).
- **Data:** *IC*:  $u(x, 0) = -\sin(\pi x)$ . *BC*:  $u(-1, t) = u(1, t) = 0$ .
- **Physical terms:**  $u_t$  (temporal change),  $u, u_x$  (self-advection),  $\nu u_{xx}$  (viscous smoothing).
- **Scaling:** With  $x = L\hat{x}$ ,  $t = T\hat{t}$ ,  $u = U\hat{u}$  and  $T=L/U$ , equation becomes  $\hat{u}_{\hat{t}} + \hat{u}\widehat{\hat{u}_{\hat{x}}} = \text{Re}^{-1}\widehat{\hat{u}_{\hat{x}\hat{x}}}$ . We're already using a dimensionless  $\nu = 0.01/\pi$ .

$$u_t + u u_x = \nu u_{xx}, \quad (x, t) \in [-1, 1] \times [0, 1].$$
$$u(x, 0) = -\sin(\pi x), \quad u(-1, t) = 0, \quad u(1, t) = 0.$$

$$x = L\hat{x}, \quad t = T\hat{t}, \quad u = U\hat{u}, \quad T = \frac{L}{U} \Rightarrow \hat{u}_{\hat{t}} + \hat{u}\hat{u}_{\hat{x}} = \text{Re}^{-1}\hat{u}_{\hat{x}\hat{x}},$$

$$\text{where,} \quad \text{Re} = \frac{UL}{\nu}.$$

# Function spaces, well-posedness, and trial functions (soft vs hard constraints)

---

- **Function space:** For Dirichlet BC, natural space  $V = \{v \in H^1(\Omega) : v(\pm 1, t) = 0\}$ .
- **Well-posedness (viscous):** Existence/uniqueness; internal shock of inviscid case becomes a viscous layer of width  $O(\nu)$ .
- **PINN trial:** Smooth parametric family  $u_\theta$  (*MLP with tanh*) dense in  $C(K)$ .
- **BC/IC handling:**
  - **Soft constraints:** penalties in loss (used here; simple and flexible).
  - **Hard constraints:** encode IC/BC in the **ansatz**, e.g., multiply by factors that vanish on BC set; reduces penalty tuning but can complicate architecture.

$$V := \{v \in H^1(\Omega) : v(\pm 1, t) = 0\}, \quad u : \Omega \rightarrow \mathbb{R} \text{ solves } u_t + uu_x - \nu u_{xx} = 0.$$
$$u_\theta(x, t) = \text{MLP}_\theta([x, t]), \quad \theta \in \mathbb{R}^p, \quad (\text{soft BC/IC via penalties}).$$

Hard-BC trial (example):  $u_\theta(x, t) = (1 - x^2) t \tilde{u}_\theta(x, t) \Rightarrow u(\pm 1, t) = 0, u(x, 0) = 0$  by construction.

# Strong residual and continuous PINN objective; normalization & chain rule

---

- **Residual operator:**  $\mathcal{N}[u] = u_t + u u_x - \nu u_{xx}$ .
- **Continuous objective:** Area integral of residual squared (physics) + lower-dimensional IC/BC integrals; weights balance terms.
- **Normalization:** Inputs scaled to  $[-1,1]$  for stable optimization; chain rule restores physical derivatives automatically in autodiff.

$$\mathcal{N}[u] := u_t + u u_x - \nu u_{xx}, \quad r_\theta(x, t) := \mathcal{N}[u_\theta](x, t).$$

$$\boxed{\mathcal{J}(\theta) = \lambda_f \int_{\Omega} |r_\theta|^2 \rho dx dt + \lambda_{ic} \int_{-1}^1 |u_\theta(x, 0) - u_0(x)|^2 \rho_{ic} dx + \lambda_{bc} \int_0^1 (|u_\theta(-1, t)|^2 + |u_\theta(1, t)|^2) \rho_{bc} dt}$$

$$\hat{x} = \frac{2(x - X_{min})}{X_{max} - X_{min}} - 1, \quad \hat{t} = \frac{2(t - T_{min})}{T_{max} - T_{min}} - 1, \quad u_x = u_{\hat{x}} \alpha_x, \quad u_{xx} = u_{\hat{x}\hat{x}} \alpha_x^2, \quad u_t = u_{\hat{t}} \alpha_t.$$

# Discrete Monte-Carlo estimator, sampling, and variance control

---

- **Random collocation:** Approximate integrals with sample means; resample each epoch → unbiased stochastic gradients.
- **Counts:** Typical ratio  $N_f:N_{ic}:N_{bc} \approx 10:1:1$  (adjust for problem difficulty)
- **Variance reduction:** Batch the interior, **stratify** near internal layer, or use **residual-based adaptive refinement (RAR)**: sample more where  $|r_\theta|$  is large.
- **Importance sampling:** If sampling with density  $q$ , weight terms by  $\rho/q$  to keep estimator unbiased.

$$\hat{\mathcal{J}}(\theta) = \lambda_f \frac{1}{N_f} \sum_{i=1}^{N_f} |r_\theta(x_i, t_i)|^2 + \lambda_{ic} \frac{1}{N_{ic}} \sum_{j=1}^{N_{ic}} |u_\theta(x_j^{ic}, 0) - u_0(x_j^{ic})|^2 + \lambda_{bc} \frac{1}{N_{bc}} \sum_{k=1}^{N_{bc}} (|u_\theta(-1, t_k^{bc})|^2 + |u_\theta(1, t_k^{bc})|^2).$$

$$\text{Var}\left[\frac{1}{N_f} \sum r_\theta^2\right] = \frac{1}{N_f} \text{Var}(r_\theta^2) \quad \Rightarrow \quad \text{increase } N_f, \text{ stratify, or use } q \propto |r_\theta| \text{ (RAR).}$$

# Autodiff construction and gradients; architecture & initialization

---

- **Autodiff pipeline:** Evaluate  $u_\theta(x, t) \rightarrow$  get  $u_t, u_x, u_{xx}$  via nested autograd.grad with create\_graph=True → form  $r_\theta$ .
- **SGD/Adam gradient:** Sum of residual-weighted sensitivity terms plus IC/BC sensitivities.
- **Architecture:** MLP with tanh (smooth derivatives); width/depth tuned to feature complexity (e.g., width 50–64, depth 8–10).
- **Initialization:** Xavier for linears; tanh avoids exploding derivatives; for high-freq content, consider **Fourier features** or **SIREN** (sine activations).

$$r_\theta(x, t) = \partial_t u_\theta + u_\theta \partial_x u_\theta - \nu \partial_{xx} u_\theta.$$

$$\nabla_\theta \hat{\mathcal{J}} = \frac{2\lambda_f}{N_f} \sum_i r_\theta(x_i, t_i) \nabla_\theta r_\theta(x_i, t_i) + \frac{2\lambda_{ic}}{N_{ic}} \sum_j \Delta_j^{ic} \nabla_\theta u_\theta(x_j^{ic}, 0) + \frac{2\lambda_{bc}}{N_{bc}} \sum_k \sum_{s \in \{-1, 1\}} u_\theta(s, t_k^{bc}) \nabla_\theta u_\theta(s, t_k^{bc}),$$

with  $\Delta_j^{ic} := u_\theta(x_j^{ic}, 0) - u_0(x_j^{ic})$ .

$$\phi_{Fourier}(z) = [\sin(2\pi Bz), \cos(2\pi Bz)], \quad \text{SIREN: } u_\theta(z) = \sin(\omega_0 W_L \cdots \sin(\omega_0 W_1 z)).$$

# Optimization schedule, weights, and practical stabilizers

---

- **Two-phase weights:** IC/BC heavy early → enforce constraints; later increase physics weight to refine interior.
- **Adam → L-BFGS:** Adam explores with fresh samples; L-BFGS polishes with **fixed samples in closure** (deterministic line search).
- **Stabilizers:** gradient clipping; batched interior loss; learning-rate schedule; early stopping by residual plateau; checkpointing.
- **Stopping:** Residual mean stabilizes, IC/BC RMSE below tolerance, and relative  $L^2$  vs reference stops improving.

$$(\lambda_f, \lambda_{ic}, \lambda_{bc}) = \begin{cases} (1, 200, 200), & \text{early,} \\ (2, 100, 100), & \text{late.} \end{cases} \quad \theta_{k+1}^{Adam} = \theta_k - \eta \frac{\hat{m}_k}{\sqrt{\hat{v}_k} + \epsilon}.$$

L-BFGS:  $p_k = -H_k \nabla \hat{\mathcal{J}}(\theta_k)$ ,  $\theta_{k+1} = \theta_k + \alpha_k p_k$ ,  $\alpha_k$  satisfies strong Wolfe;  $H_k$  via two-loop recursion.



# Diagnostics: residual stats, IC/BC RMSE, energy law

- **Residual stats:** mean/median/max of  $|r_\theta|$  on a holdout grid (physics fidelity).
- **IC/BC RMSE:** errors at  $t = 0$  and  $x = \pm 1$  (constraint satisfaction).
- **Energy decay:**  $E(t) = \frac{1}{2} \int_{-1}^1 u^2 dx$  should be non-increasing for viscous flows.

Residual stats on holdout:  $\mu = \text{mean}(|r_\theta|)$ ,  $\text{med} = \text{median}(|r_\theta|)$ ,  $M = \max(|r_\theta|)$ .

$$\text{RMSE}_{IC} = \sqrt{\frac{1}{n_x} \sum_i (u_\theta(x_i, 0) - u_0(x_i))^2}, \quad \text{RMSE}_{BC, \pm} = \sqrt{\frac{1}{n_t} \sum_n u_\theta(\pm 1, t_n)^2}.$$

$$\frac{d}{dt} \frac{1}{2} \int_{-1}^1 u^2 dx = -\nu \int_{-1}^1 u_x^2 dx \leq 0 \quad (\text{by parts; } u(\pm 1, t) = 0).$$



# Benchmark, error analysis, and interpreting discrepancies

---



- **Reference method:** Explicit FD — upwind advection, centered diffusion, forward Euler; adaptive  $\Delta t$  for CFL/diffusion stability.
- **Global metric:** relative  $L^2$  on the space–time grid.
- **Localized error band:** Internal layer near  $x \approx 0$  after  $t \gtrsim 0.3$  → small phase/width mismatches create a narrow bright stripe in  $|u_\theta - u^{FD}|$  while global  $L^2$  stays small.

$$(u_x)_i \approx \begin{cases} \frac{u_i - u_{i-1}}{\Delta x}, & u_i \geq 0, \\ \frac{u_{i+1} - u_i}{\Delta x}, & u_i < 0, \end{cases} \quad (u_{xx})_i \approx \frac{u_{i+1} - 2u_i + u_{i-1}}{\Delta x^2}, \quad u_i^{n+1} = u_i^n - \Delta t u_i^n (u_x)_i^n + \nu \Delta t (u_{xx})_i^n.$$

$$\varepsilon_{rel} = \frac{\|u_\theta - u^{FD}\|_{L^2(\Omega)}}{\|u^{FD}\|_{L^2(\Omega)}} = \frac{\left(\int_{\Omega} |u_\theta - u^{FD}|^2\right)^{1/2}}{\left(\int_{\Omega} |u^{FD}|^2\right)^{1/2}}.$$

$$u(x, t) \approx U\left(\frac{x - x_\star(t)}{\nu}\right) \Rightarrow \delta x_\star = O(\nu) \implies |u_\theta - u| = O(1) \text{ on width } O(\nu).$$

# Benchmark, error analysis, and interpreting discrepancies



# Extensions and best practices

---

- **Weak-form PINNs:** Integrate diffusion by parts; use first derivatives only; better conditioning for stiff/high-order operators.
- **Adaptive sampling (RAR):** Replenish collocation with points of high  $|r_\theta|$ ; or stratify hard subregions.
- **Fourier features / SIREN / positional encodings:** Combat spectral bias and resolve thin layers.
- **Domain/time decomposition (XPINNs):** Subdomains with interface conditions; parallel training; local refinement.
- **Augmented Lagrangian for IC/BC:** Reduce manual weight tuning; enforce constraints more tightly.
- **Inverse problems:** Make parameters (e.g.,  $\nu$ ) trainable; add sparse sensor data to loss.

$$\text{Weak form: } \int_{\Omega} (u_t + uu_x) v \, dx \, dt + \nu \int_{\Omega} u_x v_x \, dx \, dt = 0, \quad \forall v \in V.$$

RAR sampling density:  $q(x, t) \propto |r_\theta(x, t)| \Rightarrow$  importance weights  $\frac{\rho}{q}$  keep estimator unbiased.

$$\nu = \text{softplus}(\phi), \quad \mathcal{J}(\theta, \phi) = \lambda_f \|\mathcal{N}_{\nu(\phi)}[u_\theta]\|_{L^2(\Omega)}^2 + \lambda_{ic} \cdots + \lambda_{bc} \cdots + \lambda_{data} \sum_m |u_\theta(x_m, t_m) - y_m|^2.$$